

# Final project

July 7, 2019

```
In [44]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.model_selection import cross_validate
from sklearn.svm import SVC
import seaborn as sns
from tqdm import tqdm_notebook
from sklearn.preprocessing import OneHotEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import roc_auc_score
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report

sns.set();

%matplotlib inline

In [2]: df = pd.read_csv("2015StreetTreesCensus_TREES.csv")
df.head()
```

```
Out[2]:
```

	created_at	tree_id	block_id	\		the_geom	tree_dbh	stump_diam	\
0	08/27/2015	180683	348711		0	POINT (-73.84421521958048 40.723091773924274)	3	0	
1	09/03/2015	200540	315986		1	POINT (-73.81867945834878 40.79411066708779)	21	0	
2	09/05/2015	204026	218365		2	POINT (-73.93660770459083 40.717580740099116)	3	0	
3	09/05/2015	204337	217969		3	POINT (-73.93445615919741 40.713537494833226)	10	0	
4	08/30/2015	189565	223043		4	POINT (-73.97597938483258 40.66677775537875)	21	0	

	curb_loc	status	health		spc_latin	...	\
0	OnCurb	Alive	Fair		Acer rubrum	...	
1	OnCurb	Alive	Fair		Quercus palustris	...	
2	OnCurb	Alive	Good	Gleditsia triacanthos var.	inermis	...	
3	OnCurb	Alive	Good	Gleditsia triacanthos var.	inermis	...	
4	OnCurb	Alive	Good		Tilia americana	...	

	st_assem	st_senate	nta	nta_name	boro_ct	state	Latitude	\
0	28	16	QN17	Forest Hills	4073900	New York	40.723092	
1	27	11	QN49	Whitestone	4097300	New York	40.794111	
2	50	18	BK90	East Williamsburg	3044900	New York	40.717581	
3	53	18	BK90	East Williamsburg	3044900	New York	40.713537	
4	44	21	BK37	Park Slope-Gowanus	3016500	New York	40.666778	

	longitude	x_sp	y_sp
0	-73.844215	1.027431e+06	202756.768749
1	-73.818679	1.034456e+06	228644.837379
2	-73.936608	1.001823e+06	200716.891267
3	-73.934456	1.002420e+06	199244.253136
4	-73.975979	9.909138e+05	182202.425999

[5 rows x 42 columns]

In [146]: df.shape

Out[146]: (683788, 42)

In [147]: df.dropna().shape

Out[147]: (652118, 42)

In [3]: df.columns

Out[3]: Index(['created\_at', 'tree\_id', 'block\_id', 'the\_geom', 'tree\_dbh',  
'stump\_diam', 'curb\_loc', 'status', 'health', 'spc\_latin', 'spc\_common',  
'steward', 'guards', 'sidewalk', 'user\_type', 'problems', 'root\_stone',  
'root\_grate', 'root\_other', 'trnk\_wire', 'trnk\_light', 'trnk\_other',  
'brnch\_ligh', 'brnch\_shoe', 'brnch\_othe', 'address', 'zipcode',  
'zip\_city', 'cb\_num', 'borocode', 'boroname', 'cnclldist', 'st\_assem',  
'st\_senate', 'nta', 'nta\_name', 'boro\_ct', 'state', 'Latitude',  
'longitude', 'x\_sp', 'y\_sp'],  
dtype='object')

In [4]: df.shape

Out[4]: (683788, 42)

In [196]: class eda:  
def \_\_init\_\_(self, data):

```

        self.df = data

    def plot(self, col1, col2, block):
        df = self.df
        # plt.scatter(df[col1], df[col2], c = df['boroname'])
        # plt.xlabel(col1)
        # plt.ylabel(col2)
        # plt.title("%s vs %s" %(col1,col2))
        sns.lmplot(col1, col2, data=df, hue=block, fit_reg=False, scatter_kws={"s":
    def plot_noblock(self, new_df, col1, col2):
        sns.lmplot(col1, col2, data=new_df, fit_reg=False)

    def boxplot(self, col1, col2):
        df = self.df
        sns.boxplot(x=col1, y=col2, data= df)

    def hist(self, col):
        plt.hist(df[col])


    def check_corr(self, col):
        return self[col].corr()

```


```


In [6]: ed = eda(df)
        ed.plot('longitude', 'Latitude', 'boroname')


```




```
In [131]: for name in df.boroname.unique():
           df_temp = df.loc[df.boroname == name, :].loc[df.status != 'Stump', :]
           eda_temp = eda(df_temp)
           eda_temp.plot('longitude', 'Latitude', 'status')
```









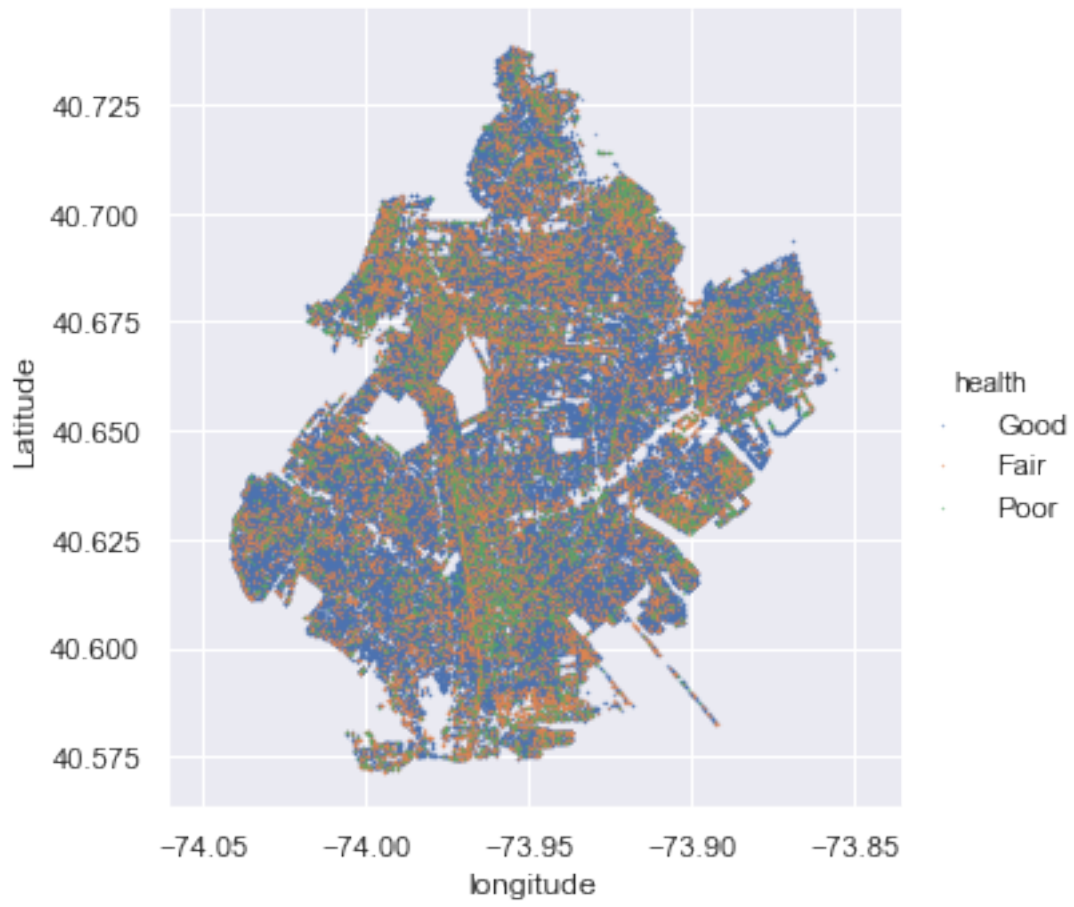









```
In [126]: for name in df.boriname.unique():  
           df_temp = df.loc[df.boriname == name,:]  
           eda_temp = eda(df_temp)  
           eda_temp.plot('longitude', 'Latitude', 'health')
```












```
In [110]: ed.plot('longitude', 'Latitude', 'status')
```



```
In [116]: #ed.plot_noblock(df.loc[df.status=='Alive'], 'longitude', 'Latitude')
```

```
In [45]: svm = SVC(kernel = 'poly')
log = LogisticRegression()
knn = KNeighborsClassifier(n_neighbors=5)
```

```
def train_test(model, X, y, multiclass = False):
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_s
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    y_prob = model.predict_proba(X_test)
    if multiclass == False:
        return(np.mean(y_test!=y_pred), roc_auc_score(y_test, y_prob[:,1]))
    else:
        return(classification_report(y_test, y_pred))
```

```
#logistic regression on geolocation only
```

```
train_test(log, df.loc[df.status != 'Stump', ['longitude', 'Latitude']],
           df.status.loc[df.status != 'Stump'])
```

```
/Users/royluo/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning
```

```
Out[45]: (0.020739039383908666, 0.5169118446337606)
```

```
In [190]: df1 = df.dropna().loc[df.status == 'Alive']
```

```
In [195]: def train_test_block(model, X, y, block_col, multiclass = False):
    for region in X[block_col].unique():
        X_block = X.loc[X[block_col] == region, :]
        X_block = X_block.drop(columns = block_col)
        y_block = y.loc[X[block_col] == region]
        X_train, X_test, y_train, y_test = train_test_split(X_block, y_block, test_s
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)
        y_prob = model.predict_proba(X_test)
        if multiclass == False:
            print(region, np.mean(y_test!=y_pred), roc_auc_score(y_test, y_prob[:,1])
        else:
            print(region, classification_report(y_test, y_pred))
```

```
X_h = df1[['longitude', 'Latitude', 'tree_dbh', 'boroname']]
#X_h = pd.concat([pd.get_dummies(df1[['steward', 'sidewalk', 'curb_loc']]),
#                df1[['longitude', 'Latitude', 'tree_dbh', 'boroname']]], axis = 1)
y_h = df1.health.loc[df1.status == 'Alive']
```

```
ada = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=2), n_estim
train_test_block(knn, X_h, y_h , 'boroname', True)
```

Queens		precision	recall	f1-score	support
--------	--	-----------	--------	----------	---------

Fair	0.31	0.17	0.22	6799
Good	0.84	0.94	0.89	38872
Poor	0.34	0.04	0.07	1922

micro avg	0.79	0.79	0.79	47593
macro avg	0.50	0.38	0.39	47593
weighted avg	0.74	0.79	0.76	47593

Brooklyn		precision	recall	f1-score	support
----------	--	-----------	--------	----------	---------

Fair	0.33	0.17	0.22	4947
Good	0.84	0.94	0.89	27713
Poor	0.38	0.05	0.08	1287



micro avg	0.80	0.80	0.80	33947
macro avg	0.51	0.39	0.40	33947
weighted avg	0.75	0.80	0.76	33947

Manhattan		precision	recall	f1-score	support
Fair	0.30	0.19	0.23	2239	
Good	0.78	0.91	0.84	9514	
Poor	0.30	0.05	0.08	732	

micro avg	0.73	0.73	0.73	12485
macro avg	0.46	0.38	0.38	12485
weighted avg	0.67	0.73	0.69	12485

Staten Island		precision	recall	f1-score	support
Fair	0.37	0.22	0.28	2927	
Good	0.85	0.94	0.89	16539	
Poor	0.36	0.05	0.09	818	

micro avg	0.80	0.80	0.80	20284
macro avg	0.53	0.41	0.42	20284
weighted avg	0.76	0.80	0.77	20284

Bronx		precision	recall	f1-score	support
Fair	0.34	0.15	0.21	2266	
Good	0.84	0.95	0.89	13269	
Poor	0.24	0.03	0.05	582	

micro avg	0.81	0.81	0.81	16117
macro avg	0.47	0.38	0.38	16117
weighted avg	0.75	0.81	0.77	16117

In [161]: train\_test\_block(ada, X\_h, y\_h , 'boroname', True)

Queens		precision	recall	f1-score	support
Fair	0.49	0.00	0.01	6799	
Good	0.82	1.00	0.90	38872	
Poor	0.50	0.00	0.01	1922	

micro avg	0.82	0.82	0.82	47593
macro avg	0.60	0.34	0.31	47593
weighted avg	0.76	0.82	0.74	47593

Brooklyn		precision	recall	f1-score	support
Fair	0.44	0.00	0.01	4947	
Good	0.82	1.00	0.90	27713	
Poor	0.17	0.00	0.00	1287	
micro avg	0.82	0.82	0.82	33947	
macro avg	0.48	0.33	0.30	33947	
weighted avg	0.74	0.82	0.74	33947	

Manhattan		precision	recall	f1-score	support
Fair	0.39	0.01	0.02	2239	
Good	0.76	1.00	0.86	9514	
Poor	0.33	0.01	0.01	732	
micro avg	0.76	0.76	0.76	12485	
macro avg	0.50	0.34	0.30	12485	
weighted avg	0.67	0.76	0.66	12485	

Staten Island		precision	recall	f1-score	support
Fair	0.45	0.01	0.01	2927	
Good	0.82	1.00	0.90	16539	
Poor	0.64	0.01	0.02	818	
micro avg	0.82	0.82	0.82	20284	
macro avg	0.64	0.34	0.31	20284	
weighted avg	0.76	0.82	0.73	20284	

Bronx		precision	recall	f1-score	support
Fair	0.35	0.00	0.01	2266	
Good	0.82	1.00	0.90	13269	
Poor	0.00	0.00	0.00	582	
micro avg	0.82	0.82	0.82	16117	
macro avg	0.39	0.33	0.30	16117	
weighted avg	0.73	0.82	0.74	16117	

```
In [194]: log = LogisticRegression(solver = 'lbfgs', multi_class = 'multinomial', class_weight =
        train_test_block(log, X_h, y_h , 'boroname', True)

/Users/royluo/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:758: ConvergenceWarning:
  "of iterations.", ConvergenceWarning)
```

Queens		precision	recall	f1-score	support
	Fair	0.17	0.10	0.12	6799
	Good	0.83	0.45	0.58	38872
	Poor	0.05	0.60	0.09	1922
	micro avg	0.40	0.40	0.40	47593
	macro avg	0.35	0.38	0.27	47593
	weighted avg	0.71	0.40	0.50	47593

/Users/royluo/anaconda3/lib/python3.6/site-packages/sklearn/linear\_model/logistic.py:758: ConvergenceWarning: "of iterations.", ConvergenceWarning)

Brooklyn		precision	recall	f1-score	support
	Fair	0.15	0.09	0.11	4947
	Good	0.84	0.45	0.59	27713
	Poor	0.05	0.67	0.10	1287
	micro avg	0.41	0.41	0.41	33947
	macro avg	0.35	0.40	0.27	33947
	weighted avg	0.71	0.41	0.50	33947

/Users/royluo/anaconda3/lib/python3.6/site-packages/sklearn/linear\_model/logistic.py:758: ConvergenceWarning: "of iterations.", ConvergenceWarning)

Manhattan		precision	recall	f1-score	support
	Fair	0.20	0.17	0.18	2239
	Good	0.82	0.46	0.59	9514
	Poor	0.09	0.62	0.15	732
	micro avg	0.41	0.41	0.41	12485
	macro avg	0.37	0.42	0.31	12485
	weighted avg	0.67	0.41	0.49	12485

/Users/royluo/anaconda3/lib/python3.6/site-packages/sklearn/linear\_model/logistic.py:758: ConvergenceWarning: "of iterations.", ConvergenceWarning)

Staten Island		precision	recall	f1-score	support
---------------	--	-----------	--------	----------	---------

Fair	0.16	0.22	0.19	2927
Good	0.85	0.33	0.47	16539
Poor	0.05	0.61	0.09	818
micro avg	0.32	0.32	0.32	20284
macro avg	0.35	0.39	0.25	20284
weighted avg	0.72	0.32	0.42	20284

```
/Users/royluo/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:758: ConvergenceWarning:
  "of iterations.", ConvergenceWarning)
```

Bronx		precision	recall	f1-score	support
	Fair	0.15	0.13	0.14	2266
	Good	0.86	0.39	0.53	13269
	Poor	0.05	0.68	0.09	582
	micro avg	0.36	0.36	0.36	16117
	macro avg	0.35	0.40	0.25	16117
	weighted avg	0.73	0.36	0.46	16117

```
In [95]: #logistic regression on geolocation and tree diameter
train_test(log, df.loc[df.status != 'Stump', ['longitude', 'Latitude', 'tree_dbh']],
           df.status.loc[df.status != 'Stump'])
```

```
/Users/royluo/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning:
  FutureWarning)
```

```
Out[95]: (0.020739039383908666, 0.7695517831463828)
```

```
In [92]: #knn based on only geolocation
train_test(knn, df.loc[df.status != 'Stump', ['longitude', 'Latitude']],
           df.status.loc[df.status != 'Stump'])
```

```
Out[92]: (0.021392060167983968, 0.6251759179474983)
```

```
In [103]: #knn based on geolocation and tree diameter
#notice there isn't much improvement
train_test(knn, df.loc[df.status != 'Stump', ['longitude', 'Latitude', 'tree_dbh']],
           df.status.loc[df.status != 'Stump'])
```

```
Out[103]: (0.021692299608938127, 0.6389132389055914)
```

```
In [90]: # testing best k parameter based on only geolocation
for k in tqdm_notebook(range(20)):
    knn_k = KNeighborsClassifier(n_neighbors=k+1)

    print([k+1, train_test(knn_k, df.loc[df.status != 'Stump', ['longitude', 'Latitude',
    df.status.loc[df.status != 'Stump']]))
```

```
HBox(children=(IntProgress(value=0, max=20), HTML(value='')))
```

```
[1, (0.035600891711139636, 0.5687548361588286)]
[2, (0.021279470377626158, 0.5958590998260858)]
[3, (0.023035871107208, 0.6122720902241869)]
[4, (0.020956712978600432, 0.6192316162825082)]
[5, (0.021392060167983968, 0.6251759179474983)]
[6, (0.020851629174266476, 0.6311245042722204)]
[7, (0.02103927882486283, 0.6357120030327346)]
[8, (0.0207015094537894, 0.6389404240937061)]
[9, (0.020806593258123354, 0.6417383193822925)]
[10, (0.020663979523670127, 0.6427023927892415)]
[11, (0.020709015439813253, 0.6450854750362436)]
[12, (0.020754051355956375, 0.6453065726330747)]
[13, (0.020754051355956375, 0.64652602091471)]
[14, (0.020716521425837107, 0.647997046272927)]
[15, (0.020716521425837107, 0.6521688844919874)]
[16, (0.0207015094537894, 0.6526093472350166)]
[17, (0.020739039383908666, 0.6541132600267862)]
[18, (0.020716521425837107, 0.6548828980995676)]
[19, (0.020724027411860958, 0.6559912378999767)]
[20, (0.020709015439813253, 0.6571263245743293)]
```

```
In [102]: #qda based on geolocation and tree diameter
qda = QuadraticDiscriminantAnalysis()
train_test(qda, df.loc[df.status != 'Stump', ['longitude', 'Latitude', 'tree_dbh']],
df.status.loc[df.status != 'Stump'])
```

```
Out[102]: (0.020739039383908666, 0.7380704226174196)
```

```
In [103]: #adaboost ensemble decision tree weak learners based geolocation and tree diameter
#best performing model so far
ada = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=1), n_estimators=100)
train_test(ada, df.loc[df.status != 'Stump', ['longitude', 'Latitude', 'tree_dbh']],
df.status.loc[df.status != 'Stump'])
```

```
Out[103]: (0.020739039383908666, 0.7730273315628547)
```

```

In [112]: y = df.status.loc[df.status != 'Stump']
          X = pd.concat([pd.get_dummies(df[['boroname', 'curb_loc']]), df[['longitude', 'Latitude', 'tree_dbh']]], axis = 1)
          X = X.loc[df.status != 'Stump', :]
          ada = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=2), n_estimators=100)
          train_test(ada, X, y)

Out[112]: (0.020739039383908666, 0.7736399768734347)

In [113]: log = LogisticRegression()
          train_test(log, X, y)

/Users/royluo/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning
    FutureWarning)

Out[113]: (0.020739039383908666, 0.7706512970193482)

In [108]: df['boroname'].value_counts()

Out[108]: Queens                250551
          Brooklyn             177293
          Staten Island        105318
          Bronx                 85203
          Manhattan             65423
          Name: boroname, dtype: int64

In [169]: #finds out the top10 species
          top10_species = df['spc_common'].value_counts()[:2].index
          df['spc_common'].value_counts()[:2]

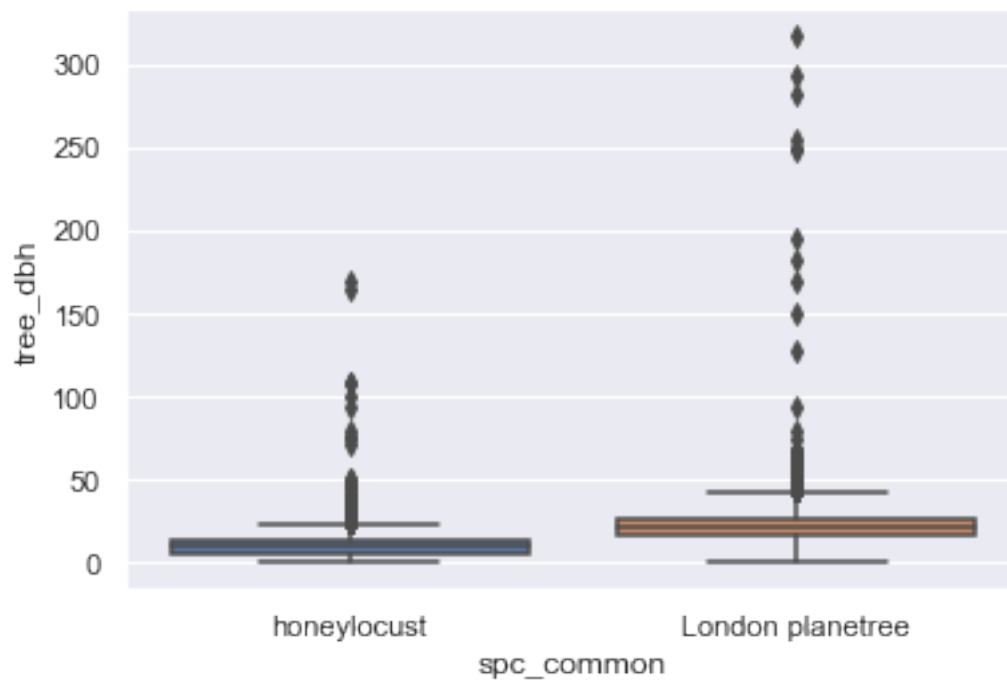
Out[169]: London planetree      87014
          honeylocust           64264
          Name: spc_common, dtype: int64

In [187]: df_10 = df.loc[df['spc_common'].apply(lambda x: True if x in top10_species else False)]
          df_10 = df_10.dropna()

In [188]: #boxplot of log(tree_dbh vs top 10 species)

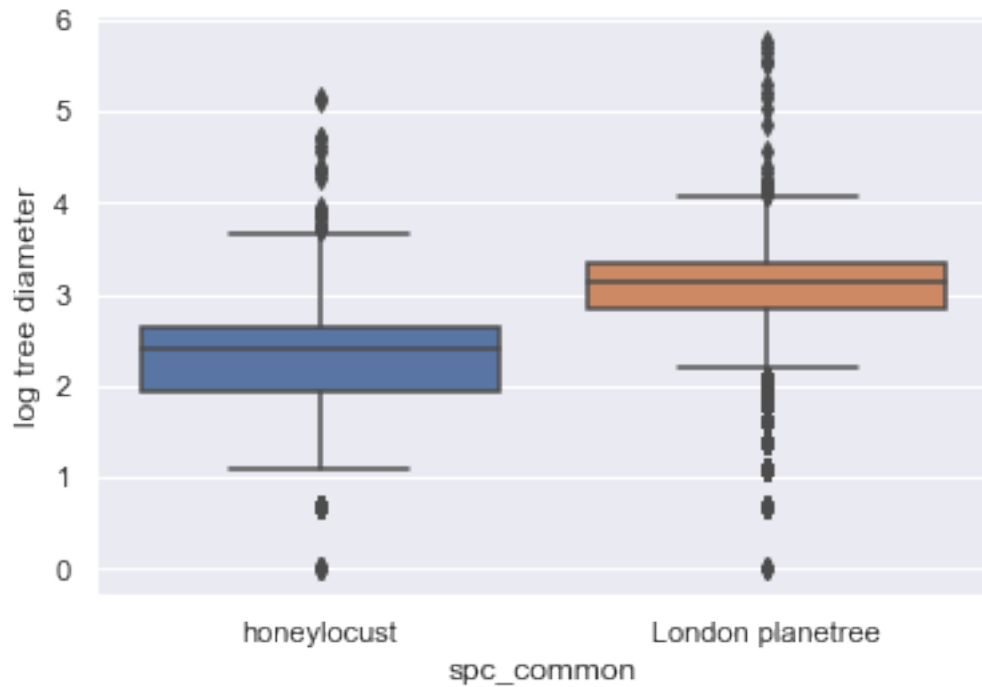
          #df_10.tree_dbh = np.log(df_10.tree_dbh)
          ed_10 = eda(df_10.loc[df_10.status != 'Stump', :])
          ed_10.boxplot('spc_common', 'tree_dbh')

```



```
In [189]: df_10.tree_dbh = np.log(df_10.tree_dbh+1)
ed_10 = eda(df_10.loc[df_10.status!= 'Stump',:])
ed_10.boxplot('spc_common', 'tree_dbh')
plt.ylabel('log tree diameter')
```

```
Out[189]: Text(0, 0.5, 'log tree diameter')
```




```
In [204]: from scipy.stats import ttest_ind
          ttest_ind(df_10.tree_dbh.loc[df_10.spc_common == "honeylocust"],
                    df_10.tree_dbh.loc[df_10.spc_common == "London planetree"],
                    equal_var = True)
```

```
Out[204]: Ttest_indResult(statistic=-258.33265678621683, pvalue=0.0)
```

```
In [173]: #geo-location of top 10 species, looks like there is clustering behavior
          #and some species only appear in certain area
          #It might be possible to predict the type of species based on other variables
          ed_10.plot('longitude', 'Latitude', 'spc_common')
```





```
In [174]: X_10 = pd.concat([pd.get_dummies(df_10[['status', 'health',
        'steward', 'sidewalk', 'boroname', 'curb_loc'])), df_10[['longitude', 'Latitude',
        'spc_common']]], axis=1)

y_10 = df_10.spc_common

In [179]: knn = KNeighborsClassifier(n_neighbors=5)
train_test(knn, X_10, y_10)

Out[179]: (0.16573015138494082, 0.9007219691600991)

In [180]: ada = AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=2), n_estimators=100)
train_test(ada, X_10, y_10)

Out[180]: (0.15859059959013685, 0.9176699487551436)

In [181]: # import os
# beep = lambda x: os.system("echo -n '\a';sleep 0.2;" * x)
# beep(3)

In [182]: log = LogisticRegression()
train_test(log, X_10, y_10)
```

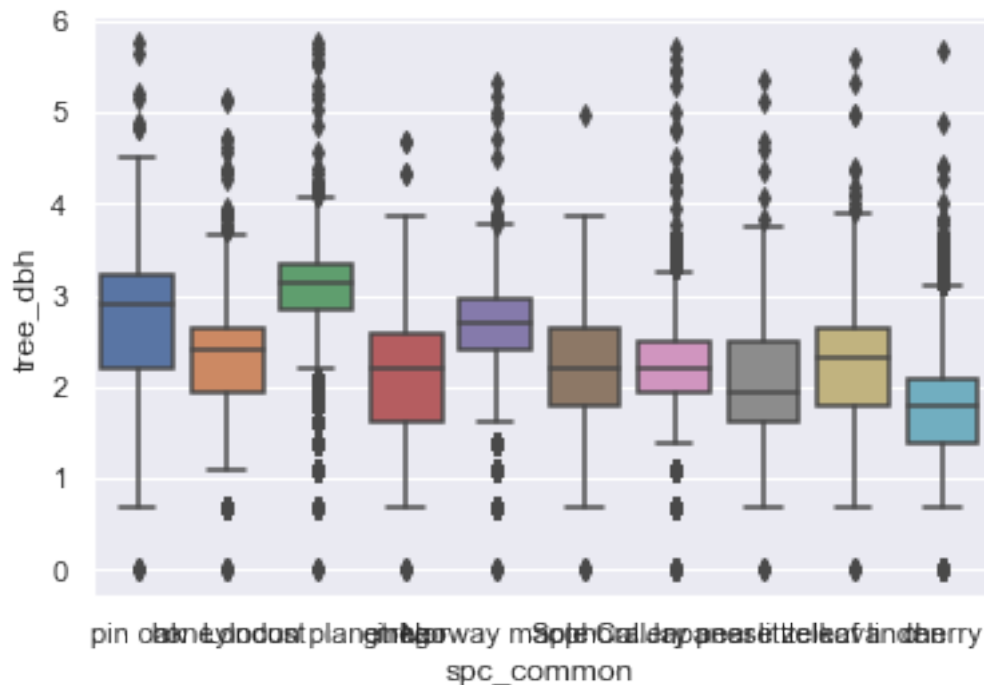
```
/Users/royluo/anaconda3/lib/python3.6/site-packages/sklearn/linear_model/logistic.py:433: FutureWarning
```

```
Out[182]: (0.21045151054406028, 0.8708595597309167)
```

```
In [200]: top10_species = df['spc_common'].value_counts()[:10].index
df_10 = df.loc[df['spc_common'].apply(lambda x: True if x in top10_species else False)]
df_10 = df_10.dropna()
df['spc_common'].value_counts()[:10]
```

```
Out[200]: London planetree      87014
honeylocust      64264
Callery pear      58931
pin oak      53185
Norway maple      34189
littleleaf linden  29742
cherry      29279
Japanese zelkova  29258
ginkgo      21024
Sophora      19338
Name: spc_common, dtype: int64
```

```
In [201]: df_10.tree_dbh = np.log(df_10.tree_dbh+1)
ed_10 = eda(df_10.loc[df_10.status!= 'Stump',:])
ed_10.boxplot('spc_common', 'tree_dbh')
```



```

In [202]: X_10 = pd.concat([pd.get_dummies(df_10[['status', 'health',
        'steward', 'sidewalk', 'boroname','curb_loc']]), df_10[['longitude','Latitude

y_10 = df_10.spc_common

train_test(knn, X_10, y_10, True)

Out[202]: '                precision    recall  f1-score   support\n\n      Callery pear

In [ ]:

```